



Vidyavardhini's College of Engineering and Technology, Vasai
Department of Computer Science & Engineering (Data Science)

AY: 2025-26

Class:	BE-CSE(DS)	Semester:	VII
Course Code:	CSDOL7011	Course Name:	NLP Lab

Name of Student:	Sahil Salunke
Roll No. :	45
Experiment No.:	9
Title of the Experiment:	Training and Evaluating a Text Classification Model Using Proper Experimental Methodology
Date of Performance:	
Date of Submission:	

Evaluation

Performance Indicator	Max. Marks	Marks Obtained
Performance	5	
Understanding	5	
Journal work and timely submission	10	
Total	20	

Performance Indicator	Exceed Expectations (EE)	Meet Expectations (ME)	Below Expectations (BE)
Performance	4-5	2-3	1
Understanding	4-5	2-3	1
Journal work and timely submission	8-10	5-8	1-4

Checked by

Name of Faculty :

Signature :

Date :



Vidyavardhini's College of Engineering and Technology, Vasai
Department of Computer Science & Engineering (Data Science)

Aim: To implement a text classification model and evaluate its performance using standard experimental procedures including data splitting, cross-validation, and evaluation metrics.

Objective: To build and evaluate a text classification model using standard machine learning methodology and evaluation metrics.

Tools Required:

1. Python (Jupyter Notebook or Google Colab)
2. scikit-learn
3. pandas, matplotlib
4. Dataset: SMS Spam Collection Dataset or any labeled text classification dataset

Procedure:

1. Import required libraries:
 - a. import pandas as pd
 - b. from sklearn.feature_extraction.text import TfidfVectorizer
 - c. from sklearn.model_selection import train_test_split, cross_val_score
 - d. from sklearn.naive_bayes import MultinomialNB
 - e. from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
 - f. import matplotlib.pyplot as plt
 - g. import seaborn as sns
2. Load the dataset:
 - a. For SMS Spam Dataset: Download from [UCI ML Repository](#)
3. Preprocess the text:

Lowercase conversion, stopwords removal (optional), and TF-IDF feature extraction



Vidyavardhini's College of Engineering and Technology, Vasai
Department of Computer Science & Engineering (Data Science)

4. Split the dataset:

Use `train_test_split()` to divide into training and testing sets (e.g., 80%-20%)

5. Train the model:

Use `MultinomialNB()` or `LogisticRegression()` classifier

6. Evaluate the model:

- a. Predict on test data
- b. Use the following evaluation metrics:
 - i. Accuracy
 - ii. Precision
 - iii. Recall
 - iv. F1-Score
 - v. Confusion Matrix

7. (Optional): Perform 5-fold cross-validation and compare with hold-out evaluation.

8. Visualize results:

Plot confusion matrix using `seaborn.heatmap()`

Description of the Experiment:

This experiment teaches students how to build and evaluate a complete text classification system using real-world data. It emphasizes experimental methodology, including the importance of data splits, model evaluation, and fair performance comparison.

Detailed Description of the NLP Technique:

1. Text Classification:

The task of assigning a category or label to a given text (e.g., spam vs. ham). It's widely used in:



Vidyavardhini's College of Engineering and Technology, Vasai
Department of Computer Science & Engineering (Data Science)

- a. Spam filtering
- b. Sentiment analysis
- c. Topic categorization

2. Pipeline Stages:

- a. Text Preprocessing & Vectorization: Convert raw text into numeric features using TF-IDF.
- b. Model Training: Use supervised machine learning algorithms like Naive Bayes or SVM.
- c. Model Evaluation: Use appropriate metrics to evaluate the model's generalization on unseen data.

3. Evaluation Metrics:

- a. Accuracy: Ratio of correctly predicted instances
- b. Precision: $\text{True Positives} / (\text{True Positives} + \text{False Positives})$ Recall: $\text{True Positives} / (\text{True Positives} + \text{False Negatives})$ F1-Score: Harmonic mean of precision and recall
- c. Confusion Matrix: Shows TP, TN, FP, FN counts

4. Best Practices in Experimental Methodology:

- a. Train-test split ensures model evaluation on unseen data.
- b. Cross-validation helps in robust performance estimation.
- c. Random seed control improves reproducibility.

OUTPUT:



Vidyavardhini's College of Engineering and Technology, Vasai

Department of Computer Science & Engineering (Data Science)

1. Import required libraries

```
[1] import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns
```

2. Load the dataset

```
[2] url = "https://archive.ics.uci.edu/ml/machine-learning-databases/00228/smsspamcollection.zip"
```

```
[3] import zipfile, requests, io
r = requests.get(url)
z = zipfile.ZipFile(io.BytesIO(r.content))
df = pd.read_csv(z.open("SMSSpamCollection"), sep='\t', names=["label", "message"])
```

```
print("Dataset shape:", df.shape)
print(df.head())
```

```
Dataset shape: (5572, 2)
  label message
0  ham Go until jurong point, crazy.. Available only ...
1  ham Ok lar... Joking wif u oni...
2  spam Free entry in 2 a wkly comp to win FA Cup fina...
3  ham U dun say so early hor... U c already then say...
4  ham Nah I don't think he goes to usf, he lives aro...
```

3. Preprocess the text

3.1 Convert labels: ham=0, spam=1

```
[4] df['label_num'] = df['label'].map({'ham': 0, 'spam': 1})
```

```
[5] df.head()
```

	label	message	label_num
0	ham	Go until jurong point, crazy.. Available only ...	0
1	ham	Ok lar... Joking wif u oni...	0
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	1
3	ham	U dun say so early hor... U c already then say...	0
4	ham	Nah I don't think he goes to usf, he lives aro...	0

3.2 TF-IDF Vectorization

```
[6] tfidf = TfidfVectorizer(stop_words='english', lowercase=True)
X = tfidf.fit_transform(df['message'])
y = df['label_num']
```



✓ 4. Split the dataset (80-20 split)

```
[7] ✓ Os X_train, X_test, y_train, y_test = train_test_split(  
X, y, test_size=0.2, random_state=42, stratify=y  
)
```

✓ 5. Train the model

```
[8] ✓ Os model = MultinomialNB()  
model.fit(X_train, y_train)
```

↗ MultinomialNB ⓘ ⓘ
MultinomialNB()

✓ 6. Evaluate the model

```
[9] ✓ Os y_pred = model.predict(X_test)  
  
print("\n--- Evaluation Metrics ---")  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

↗ --- Evaluation Metrics ---
Accuracy: 0.9704035874439462

Classification Report:
precision recall f1-score support

0 0.97 1.00 0.98 966
1 1.00 0.78 0.88 149

accuracy 0.97 1115
macro avg 0.98 0.89 0.93 1115
weighted avg 0.97 0.97 0.97 1115

✓ Confusion Matrix

```
[10] ✓ cm = confusion_matrix(y_test, y_pred)
```

```
[11] ✓ Os cm  
  
↗ array([[966, 0],  
[ 33, 116]])
```




7. Perform 5-fold cross-validation

```
[12]
✓ Os
cv_scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')
print("\n5-Fold Cross Validation Accuracy:", cv_scores)
print("Mean CV Accuracy:", cv_scores.mean())
```

5-Fold Cross Validation Accuracy: [0.97847534 0.96681614 0.96319569 0.97127469 0.97217235]
Mean CV Accuracy: 0.970386841745095

8. Visualize results

```
[13]
✓ Os
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["Ham", "Spam"], yticklabels=["Ham", "Spam"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



Conclusion:

- The results demonstrate that text preprocessing + TF-IDF + MultinomialNB is a highly effective pipeline for spam classification.
- The model achieves high accuracy and strong precision/recall balance, making it reliable for practical SMS spam filtering systems.